



Ethical values can be both objective and yet without a foundation Aeon Essays

English for Written Communication (Open University Malaysia)



Scan to open on Studocu

Ethics has no foundation

Ethical values can be both objective and knowable
– torture really is wrong – yet not need any
foundation outside themselves

by Andrew Sepielli

Andrew Sepielli is associate professor of philosophy at the University of Toronto, and a Laurance S Rockefeller visiting faculty fellow at the Princeton University Center for Human Values. He is the author of *Pragmatist Quietism: A Meta-Ethical System* (2022).

Edited by Sam Dresser

Many academic fields can be said to ‘study morality’. Of these, the philosophical sub-discipline of normative ethics studies morality in what is arguably the least alienated way. Rather than focusing on how people and societies think and talk about morality, normative ethicists try to figure out which things are, simply, morally good or bad, and why. The philosophical sub-field of meta-ethics adopts, naturally, a ‘meta-’ perspective on the kinds of enquiry that normative ethicists engage in. It asks whether there are objectively correct answers to these questions about good or bad, or whether ethics is, rather, a realm of illusion or mere opinion.

Most of my work in the past decade has been in meta-ethics. I believe that there are truths about what’s morally right and wrong. I believe that some of these truths are objective or, as they say in the literature, ‘stance-independent’. That is to say, it’s not my or our disapproval that makes torture morally wrong; torture is wrong because, to put it simply, it hurts people a lot. I believe that these objective moral truths are knowable, and that some people are better than others are at coming to know them. You can even call them ‘moral experts’ if you wish.

Of course, not everyone agrees with all of that. Some are simply confused; they conflate ‘objective’ with ‘culturally universal’ or ‘innate’ or ‘subsumable under a few exceptionless principles’ or some such. But many people’s misgivings about moral objectivity are more clear headed and deeper. In particular, I find that some

demur because they think that, for there to be moral truths, let alone objective, knowable ones, morality would have to have a kind of 'foundation' that, in their view, is nowhere to be found. Others, anxious to help, try to show that there's a firm foundation or ultimate ground for morality after all.

It's my view that both sides of this conflict are off on the wrong foot. Morality is objective, but it neither requires nor admits of a foundation. It just kind of floats there, along with the evaluative realm more generally, unsupported by anything else. Parts of it can be explained by other parts, but the entirety of the web or network of good and evil is brute. Maybe you think that's weird and even worthy of outright dismissal. I once thought the same thing. The purpose of this essay, which is based on my book *Pragmatist Quietism: A Meta-Ethical System* (2022), is to encourage you to start seeing this aspect of the world as I now see it.

The first question we should ask is: what exactly is a 'foundation'? We can get clearer on what a foundation is by querying whether a moral theory like utilitarianism might count as one. Utilitarianism says that actions are right to the extent, and only to the extent, that they promote overall wellbeing. So, is utilitarianism in the running for being a foundation for morality? Well, it certainly purports to explain a lot when it comes to right and wrong. Why give to the poor? Promotes wellbeing. Why not punch your neighbour in the face? Doesn't promote wellbeing. Should the Bank of Canada raise interest rates this quarter? Not clear, because it's not clear whether it promotes wellbeing. And so on, and so on.

Nonetheless, utilitarianism is not what I have in mind by a 'foundation'. This is not because utilitarianism is incorrect; it is because utilitarianism is a moral theory. But a foundation is not a moral theory. It's the kind of thing that's supposed to ground, or support, or justify, moral theories, and moral claims generally, without itself being a claim within the domain of morality.

Here's another way to think about it. Suppose that a moral sceptic were to declare, along with David Hume: 'You cannot rationally infer an "ought" from an "is"!' Now imagine that I replied: 'Oh yes you can! Utilitarianism is true, and so, from the fact that an action promotes overall wellbeing, you can infer that it's what you ought to do.' I suspect that our sceptic would regard this response as unsatisfactory. 'You can't show that Hume was wrong about "ought" and "is" by just wheeling in some further "ought",' she might respond. 'To show that the move from "is" to "ought" can be rational, you would need to step outside of

“ought”-discourse entirely, and provide a...a...’ ‘And provide,’ I would finish the sceptic’s sentence, ‘what I’m calling a “foundation”.’

So a moral theory doesn’t count as a foundation. What would count? Here’s a possible candidate. One thing that philosophers of language try to do is to explain why terms and concepts refer to the things in the world that they do. Many of these theories of reference invoke the relation of causal regulation – regulation of our ‘tokening’ of the concept ‘cat’ or our use of the word ‘cat’, for instance, by the comings and goings of the long-tailed housepets that like to stretch out on the windowsill. Some philosophers have applied this theory of reference to moral terms and concepts, yielding a view on which a concept like ‘good’ refers to whichever property or cluster of properties causally regulates our employment of it. Anything that then had that property(-cluster) would therefore be good. Note that our starting point here is not a claim or theory that is, intuitively speaking, within the subject-matter of ethics. Rather we began with a theory of reference – something belonging to the philosophy of language – that purports to explain how terms and concepts across the board are anchored in the world. One might say that, in doing so, we gave ethics a foundation.

Here is another theoretical move that might count as an attempt at offering a foundation for ethics. Many philosophers these days are leery about accepting the existence of objects, processes or properties that are outside the ‘natural’ order. This may seem to present a problem for ethics, because the right and the good have the feel of being supernatural, like ghosts and auras, rather than natural, like clams and carbon. But a few philosophers have suggested that this is too quick. There may be, in Philippa Foot’s words, ‘natural goodness’. Doctors speak of a well-functioning kidney, farmers of an underdeveloped calf, and nobody takes them to be dipping into the realm of, as they say, ‘woo’. And while some philosophers have expressed suspicion about so-called ‘teleological’ features like functions and ‘final ends’, others have argued that a closer look at scientific practice reveals their explanatory value. But if there is nothing problematic about goodness in the way of a heart, there should be nothing problematic about goodness in the way of a human being. On this, as it’s sometimes called, ‘neo-Aristotelian’ picture, then, ethical features are part of the natural world.

What makes a semantic account like the causal theory of reference or a metaphysical view like neo-Aristotelian naturalism a candidate for being a foundation, while a theory like utilitarianism is not? They are capable of serving as foundations for ethics because, basically, they’re *not* ethics; they’re semantics – they’re about what words and concepts mean – or they’re metaphysics,

cataloguing what sorts of things exist in the world. Utilitarianism, by contrast, is ethics, and ethics is no more capable of hoisting itself up by its own bootstraps than is anything else. I think we can go a little further, though. While a theory like utilitarianism offers a direct explanation – maybe a good one, maybe a bad one – of what is right or good or whatnot, our causal theory of reference does not. It offers a theory of what concepts and terms refer to, which has implications for which ethical claims are true, which in turn has implications for what's right or good. But ultimately, it tells you about what things mean, while a theory like utilitarianism tells you what's right. One indicator of the difference between the respective theories' explanatory roles is the difference between them in terms of what we may call 'domain generality'. Theories like 'terms refer to the features that causally regulate their usage' or 'only things posited by the successful natural sciences exist' have implications beyond ethics – into what 'cat' means, or about whether René Descartes's postulated *res cogitans* exists – while utilitarianism is solely a theory of right and wrong, and that's it.

Now, if you were to go on the website formerly known as Twitter and search for 'foundation morality' or something similar, you'd turn up many threads about God or religion. So it's worth asking: is God the kind of thing that people like me have in mind when we talk about a 'foundation'? There's much to be said on this matter, but on the face of it, no. If someone were to claim that an action is morally wrong if and only if God forbids it, I'd take this as an ordinary moral theory on a par with the claim that an action is morally wrong if and only if it fails to promote wellbeing. If utilitarianism isn't the sort of thing that's even eligible to be a foundation, then neither is this simple version of divine command theory. Now, to be sure, there are ways of beefing up divine command theory so that it might properly be regarded as a stab at a foundation – bringing in the metaphysics of 'God's nature', for example. (It should be said: there are parallel ways of beefing up other normative ethical theories, too.) The only point I wish to make now is that 'God commands X' no more takes us 'outside of ethics' than 'X maximises overall wellbeing'. The moral relevance of each one is up for dispute, and that dispute would take place in the arena of regular old first-order moral thinking, with the rest of the normative-ethical gladiators.

So why is it so often thought that morality requires a foundation? It may seem difficult to explain a way of thinking that strikes one as so obviously correct. I, however, do not think it is correct, let alone obviously correct, and so let me try my hand. Basically, I suspect that many people think morality needs a foundation because they in some way or other assimilate the enquiry that gets called 'normative ethics' to ordinary factual enquiry, in which there do indeed seem to

be foundations/explanations for the most argued-over claims. Whether or not you accept highfalutin philosophical positions like the principle of sufficient reason, my guess is that that you would look askance at someone who said that it's going to snow tomorrow but then claimed that there was no explanation for that – that it's just a brute fact. But if that claim strains credulity, then the view in which ethics as a whole 'just floats there', as I put it, untethered from anything that might serve to explain it, is apt to strike you as downright absurd.

Correlatively, the fundamental reason why I don't think that morality requires a foundation is that I deny that the relevant sorts of ethical disputes are akin to ordinary factual disputes. They have features that make it easy to be *fooled* into thinking otherwise, but in fact they're crucially different. More specifically, disputes that get called 'normative ethics' are most like disputes that many people have labelled 'merely verbal' or 'non-substantive'. A classic example comes from William James's book *Pragmatism* (1907). A man is chasing a squirrel around a tree. Is the man thereby going around the squirrel? One disputant says 'no', because the man is always behind the squirrel. Another says 'yes', because the man is first north of the squirrel, then west, then, south, then east of it. The people in this dispute have different beliefs, to be sure; their conflict is not a conflict of desires or emotions. Still, there's a clear sense in which they're not really representing the world in different ways. The side you take in this dispute does not determine, either directly or indirectly by way of inference, the way you think any aspect of the world looks, smells, sounds, etc; nor would taking one side or the other of this dispute guide you to act in a way that achieves your aims, whatever these aims may be and whatever your powers may be. The belief, in other words, doesn't function in the way a representation like a map does.

I think the debates that tend to get called 'normative ethical' are a lot like this. The way that the world will look, smell, sound, etc if utilitarianism is true is just the way it will look, smell, sound, etc if utilitarianism is false. Taking sides for or against utilitarianism does not help us to further our ultimate goals, whatever they may happen to be, in the way that a map does. Rather, it simply changes what our ultimate goals are.

With that said, there are also some important differences between the 'utilitarianism' and 'squirrel' debates. I said that we sometimes call disputes like that about the squirrel 'merely verbal' or 'non-substantive'. We also sometimes say of them something like: 'You could say this, or you could say that. What's the point?' This is because not only is there no representational accuracy up for grabs in these debates – nothing of value seems to be afforded by them. They seem to

be, again, pointless. Not so the majority of our debates about morality and politics. This is because such debates bear on our own and others' motivations, as well as on praise, blame, esteem and so forth in a way that debates like 'squirrel' seem not to. We might say that they are significant, but not substantive. Unlike 'squirrel', they matter. But then unlike ordinary factual disputes, the *way* that they matter is not by affording accurate representation of the world.

It's these connections with motivation and emotion that fool us into assimilating disputes about utilitarianism, or the 'trolley problem', or distributive justice, to ordinary factual disputes. Because they bear on what we do and how we feel, we do not reckon that we can simply 'go either way' on them in a willy-nilly fashion. We do not regard them as arbitrary, in other words, in the way that we regard 'squirrel'. Nor do we think it's acceptable to settle them by conceptual fiat, as we would settle disputes like 'squirrel'.

Here is what I mean by that. Were I to find myself embroiled in a discussion about whether the man is going around the squirrel, I would probably try to put a stop to it by saying: 'Look, all I mean by "going around" is this...'. By contrast, suppose we were embroiled in a dispute about whether the media would be right to mothball a story in an attempt to ensure that a disfavoured candidate is not elected. Here I would not try to settle the dispute by saying, eg, 'Look, by "right", all I mean is "maximises overall wellbeing"...'. I'd see such a dispute as to be settled by argument, not by stipulation. And again, I think we can chalk up this difference to the fact that normative-ethical disputes, despite failing to afford representational significance just like 'squirrel', are significant in practical and affective ways that 'squirrel' is not.

This all puts 'normative-ethical' disputes in a strange category, and makes it difficult to know what to say about them in terms of philosophical theory. I actually consider this an advantage, for it is manifestly not obvious what to say about truth and objectivity and knowledge when it comes to ethics! This is witnessed by the fact that some super-smart philosophers think that there are objective truths about ethics, some think ethics is bullshit along the lines of alchemy, some think ethical disputes are really conflicts of desire-like attitudes in disguise, and so on. Anyone who thought ethical disputes work in such a way that one theoretical interpretation is just utterly obvious and natural and easy to state would then owe us an explanation of how so many smart people could be getting it so terribly wrong at this late stage in intellectual history.

And so, acknowledging that it is by no means obvious, here is my own theoretical interpretation. The reason why ethics neither requires nor admits of a foundation outside of itself is that, like 'squirrel' but unlike any ordinary factual disputes, the relevant kinds of ethical dispute are non-representational or, as I prefer to put it, fail to afford 'representational value'. That is to say, one does not represent or mirror or copy the world in any robust sense that is worth caring about by coming to any conclusion rather than another pursuant to such a dispute. But the sorts of extra-ethical considerations drawn from metaphysics, semantics and so on that people typically call upon to serve as 'foundations' could be relevant to ethics only by bearing on which moral beliefs, if any, were good or bad in representational respects. They're not ethically important in the ways that happiness, freedom, equality, dignity and other such things are. But since representational value and disvalue aren't on the cards when it comes to normative-ethical disputes, these considerations regarding the metaphysics of moral properties, the sense and reference of moral terms and so on, are irrelevant to fundamental ethics.

And so it would be a mistake to think, with so-called 'error theorists' or 'nihilists' about morality, that there are no such things as moral properties in the world, and so all attributions of rightness or wrongness are false. The world doesn't have to have these little moral doodads for things to be right or wrong; there just has to be happiness and unhappiness, freedom and tyranny, and so forth. It would be a mistake to think, with Elizabeth Anscombe in her influential paper 'Modern Moral Philosophy' (1958), that the moral 'ought' lacks sense, as it were, and so there is nothing that we morally ought to do.

Whether something 'lacks sense' is a semantic matter, and semantics does not bear on normative ethics. It would bear on ethics only if it went towards determining the representational values of beliefs about ethics, but there are no such values at stake. As I said at the outset, my quarrel is not only with the sceptics. Someone who attempted to wring some positive moral conclusions out of claims in semantics (eg, about the sense or reference of moral terms) or metaphysics (eg, about what would best accomplish the reduction of morality to some cluster of suitably 'natural' properties) is making the same sort of basic error. They are treating normative-ethical enquiry as representational, even though it is not.

But if neither side of a normative-ethical dispute is representing or 'mirroring' the world any more successfully than the other is, then why can't we 'go either way',

as it seems we can in 'squirrel'? How can there be a truth of the matter, if there's no possibility of accurate or inaccurate representation in any robust sense?

My basic answer is that what gives these normative-ethical debates the appearance of mattering – their conclusions' influence on motivation and affect – also makes it the case that they actually matter. There's value and significance up for grabs in these ethical disputes, then, but it's not value that inheres in representing the world in a robust sense. It's what I call 'specifically ethical value' – the value of doing the right thing for the right reason. And it's from this sort of value that I try to wring a kind of truth or correctness that's proprietary to ethics. Imagine a kind of advisor who's ideal in all non-moral respects – true beliefs about non-evaluative matters, perfect inferential abilities, etc. If we plug a particular moral belief into such a person, and she advises you to do all and only right actions, then that belief counts as true in this proprietarily ethical sense, even though the belief does not 'picture' or 'mirror' the world.

Note that my brief for ethical truth bottoms out in claims about 'specifically ethical' value, and that my argument for the irrelevance of metaphysics, semantics, etc to ethics bottoms out in claims about what I called 'representational' value. This might strike you as begging the question against the sceptic about evaluative truth and knowledge – in other words, as assuming at the outset just what I intend to demonstrate to such a sceptic. My rejoinder: yes, I do beg the question, but this, in itself, does not put me in bad company. Everyone who ventures a positive claim about some subject matter – the external world, induction, mathematical knowledge, what-have-you – rather than withholding judgment entirely, must at some point confront the so-called 'Agrippan trilemma': either posit certain facts as unexplained, or beg the question, or accept an infinite regress. If these are problems, they're not problems for me specifically; they're problems for anyone who thinks things.

So I say that the true sin lies not in question-begging, but in failing to subsume aspects of the world within a more general vindicatory framework. For example, a theory of *a priori* knowledge that explains how knowledge of that very theory is possible might beg the question, but so long as it accounts for *a priori* knowledge in general – eg, of mathematics, logic and morality – and not just *a priori* knowledge of itself, it needn't be problematic. A theory of accurate mental representation of the world that explains how our beliefs in that very theory accurately represent the world also begs the question, but this should not worry us insofar as it explains accurate mental representation across the board. These theories earn their keep by making sense of what would otherwise remain

mysterious, and so it should not trouble us if they end up vindicating themselves in the process.

I propose to attain a similar sort of explanatory unity by vindicating all claims and domains that are worthy of it – not just ethics, but everything from biochemistry to sports prognostication – fundamentally in terms of values, be these representational, specifically ethical, or other sorts of values. It is this values-first re-imagining of enquiry for which I reserve the label ‘pragmatism’. Pragmatism offers a way of making sense of ethical truth, objectivity and knowledge by ensconcing these within a more comprehensive world picture, but not in such a way that would count as providing a foundation for ethics in some allegedly more fundamental area of enquiry. What emerges is a free-floating evaluative sphere, coupled with an account of why this is not so odd or mysterious after all.

aeon.co

24 November 2023